

How Far Does a Litter Detector Travel?

CROSS-DOMAIN URBAN WASTE DETECTION ACROSS HANDHELD, DASHCAM AND DRONE IMAGERY

Jairzinho Santos

Universidad Nacional de Ingeniería
MSc in Artificial Intelligence
Computer Vision, Prof. Elian Laura Riveros

The question

Litter detection benchmarks are evaluated **in-domain**, yet any city deploying mixed sensors — citizen photos, garbage-truck dashcams, drones implicitly assumes detectors transfer between viewpoints. **Nobody had measured whether they do.** We train and evaluate across three public domains under one audited protocol:

Dataset	Viewpoint	Images	Median obj.	% small
TACO	handheld	1,500	171 px	25
RoLID-11K	dashcam	11,564	22 px	84
UAVVaste	drone	772	72 px	66

19 experiments, 4 model families (Faster R-CNN, YOLOv11n, FCN, VGG+CAM), every metric from the same pycocotools contract against frozen ground truths tables and figures are generated, never transcribed.


82%

cross-domain mAP₅₀ drop

+21%


AP₅₀ inflation from split leakage

Detectors do not travel




In-domain mean 0.639 AP₅₀; cross-domain 0.118. The dashcam domain is an **island** (AP₅₀ ≤ 0.064 both directions): its 22-px median object falls below the smallest 32-px FPN anchor. Handheld→drone retains 68% scale, not appearance, dominates the gap.

Cheap levers: resolution



Gains concentrate where objects are smallest: dashcam +8.3 AP₅₀ / +5.8 AP_s; handheld +3.2 / +3.9. Class granularity: excluding the marginal glass class (38 test instances, AP₅₀ 0.049) lifts the remaining five classes by +0.014 AP₅₀ on average.

Why they fail: TIDE error anatomy



Multiclass models lose up to **24.6 dAP to classification** vs. 1.7 to localization naming the material, not

Auditing before training: 7 defects found

Systematic checks (perceptual hashing, EXIF verification, split cross-referencing) on the released datasets surfaced:

- **H4 - video-level leakage in RoLID-11K's official splits:** 58.2% of test frames share a source video with training (87% intra-video visual redundancy). **Published baselines are inflated by scene memorization.**
- H1/H2 - TACO basename collisions across batches; ~37% pending EXIF rotations with boxes annotated on the rotated frame.
- H5/H6 - duplicate image between val and test (RoLID); near-duplicate pair crossing train/test (UAVVaste).
- H3/H7 - release folders do not encode capture hour; annotation frames inconsistent across datasets.

Cure: medallion pipeline (raw→bronze→silver →gold) with SHA-256 lineage, five corrections, and frozen group-aware splits (pHash groups / source video; iterative stratification puts every class within 0.1 pp of the 15% test target).


+8.3

AP₅₀ from 640→1024 px (dashcam)

16×


smaller model, within 1-4 AP₅₀ points

Split hygiene changes conclusions



On the *same* clean test set, the model trained on the leaky official split scores **0.844 vs. 0.695** for the cleanly trained one: **+21% relative** memorized scenes, not generalization. Group-aware (by-video) splitting removes the bias at zero cost.

41M vs 2.6M parameters, same splits



YOLOv11n stays within 1-4 AP₅₀ points of Faster R-CNN on single-class litter (0.695 vs 0.707 dashcam) lightweight deployment is viable. The gap widens on materials (0.157-0.192 vs 0.244-0.246). FCN: 0.577 test IoU; VGG-16 crops: 97.8% accuracy with CAM evidence on the object.

Takeaways & reproducibility

- **No universal litter model yet:** per-domain fine-tuning (or at least resolution adaptation) remains mandatory for mixed-sensor deployments.
- **Audit before you train:** one hidden leak rewrites a benchmark; video/burst datasets should ship group identifiers.
- **Detect class-agnostic, classify crops:** TIDE argues for decoupling material recognition (97.8% on crops) from detection.

Three released routes: demo with pretrained weights on any photo (~5 min, CPU) · regenerate every paper number from artifacts (~10 min) · full pipeline from public data. Code, weights and splits: **github.com** released with the course hand-in.